

title: "Practical  
Machine  
Learning  
Project"  
author: "Ajay  
Rathi" date:  
"23 August  
2015" output:  
html\_document

---

## SYNOPSIS

Source:

Velloso, E.; Bulling, A.; Gellersen, H.; Ugulino, W.; Fuks, H. Qualitative Activity Recognition of Weight Lifting Exercises. Proceedings of 4th International Conference in Cooperation with SIGCHI (Augmented Human '13) . Stuttgart, Germany: ACM SIGCHI, 2013.

Goal of this project is to "predict the manner in which trainers did the exercise."

Further report should describe:

"how you built your model" "how you used cross validation" "what you think the expected out of sample error is" "why you made the choices you did" Ultimately, the prediction model is to be run on the test data to predict the outcome of 20 different test cases.

First, though, I'll load the appropriate packages and set the seed for reproduceable results.

```
library(AppliedPredictiveModeling)

## Warning: package 'AppliedPredictiveModeling' was built under R version
## 3.2.2

library(caret)

## Warning: package 'caret' was built under R version 3.2.2

## Loading required package: lattice
## Loading required package: ggplot2

## Warning: package 'ggplot2' was built under R version 3.2.2

library(rattle)

## Warning: package 'rattle' was built under R version 3.2.2

## Loading required package: RGtk2

## Warning: package 'RGtk2' was built under R version 3.2.2

## Rattle: A free graphical interface for data mining with R.
## Version 3.5.0 Copyright (c) 2006-2015 Togaware Pty Ltd.
## Type 'rattle()' to shake, rattle, and roll your data.
```

```
library(rpart.plot)

## Warning: package 'rpart.plot' was built under R version 3.2.2

## Loading required package: rpart

library(randomForest)

## Warning: package 'randomForest' was built under R version 3.2.2

## randomForest 4.6-10
## Type rfNews() to see new features/changes/bug fixes.
```

## QUESTION

In the aforementioned study, six participants participated in a dumbbell lifting exercise five different ways. The five ways, as described in the study, were "exactly according to the specification (Class A), throwing the elbows to the front (Class B), lifting the dumbbell only halfway (Class C), lowering the dumbbell only halfway (Class D) and throwing the hips to the front (Class E). Class A corresponds to the specified execution of the exercise, while the other 4 classes correspond to common mistakes."

By processing data gathered from accelerometers on the belt, forearm, arm, and dumbbell of the participants in a machine learning algorithm, the question is can the appropriate activity quality (class A-E) be predicted?

## INPUT DATA

The first step is to import the data and to verify that the training data and the test data are identical.

```
# Download data.
#url_raw_training <-
  "https://d396qusza40orc.cloudfront.net/predmachlearn/pml-training.csv"
file_dest_training <- "pml-training.csv"
#download.file(url=url_raw_training, destfile=file_dest_training,
  method="curl")
#url_raw_testing <- "https://d396qusza40orc.cloudfront.net/predmachlearn/pml-
  testing.csv"
file_dest_testing <- "pml-testing.csv"
#download.file(url=url_raw_testing, destfile=file_dest_testing,
  method="curl")

# Import the data treating empty values as NA.
df_training <- read.csv(file_dest_training, na.strings=c("NA", ""),
  header=TRUE)
colnames_train <- colnames(df_training)
df_testing <- read.csv(file_dest_testing, na.strings=c("NA", ""), header=TRUE)
colnames_test <- colnames(df_testing)

# Verify that the column names (excluding classe and problem_id) are
```

```

identical in the training and test set.
all.equal(colnames_train[1:length(colnames_train)-1],
colnames_test[1:length(colnames_train)-1])

## [1] TRUE

```

## FEATURES

Having verified that the schema of both the training and testing sets are identical (excluding the final column representing the A-E class), I decided to eliminate both NA columns and other extraneous columns.

```

# Count the number of non-NAs in each col.
nonNAs <- function(x) {
  as.vector(apply(x, 2, function(x) length(which(!is.na(x)))))
}

# Build vector of missing data or NA columns to drop.
colcnts <- nonNAs(df_training)
drops <- c()
for (cnt in 1:length(colcnts)) {
  if (colcnts[cnt] < nrow(df_training)) {
    drops <- c(drops, colnames_train[cnt])
  }
}

# Drop NA data and the first 7 columns as they're unnecessary for predicting.
df_training <- df_training[,!(names(df_training) %in% drops)]
df_training <- df_training[,8:length(colnames(df_training))]

df_testing <- df_testing[,!(names(df_testing) %in% drops)]
df_testing <- df_testing[,8:length(colnames(df_testing))]

# Show remaining columns.
colnames(df_training)

## [1] "roll_belt"           "pitch_belt"          "yaw_belt"
## [4] "total_accel_belt"    "gyros_belt_x"        "gyros_belt_y"
## [7] "gyros_belt_z"        "accel_belt_x"         "accel_belt_y"
## [10] "accel_belt_z"        "magnet_belt_x"        "magnet_belt_y"
## [13] "magnet_belt_z"       "roll_arm"            "pitch_arm"
## [16] "yaw_arm"            "total_accel_arm"     "gyros_arm_x"
## [19] "gyros_arm_y"         "gyros_arm_z"         "accel_arm_x"
## [22] "accel_arm_y"         "accel_arm_z"         "magnet_arm_x"
## [25] "magnet_arm_y"        "magnet_arm_z"        "roll_dumbbell"
## [28] "pitch_dumbbell"      "yaw_dumbbell"        "total_accel_dumbbell"
## [31] "gyros_dumbbell_x"    "gyros_dumbbell_y"    "gyros_dumbbell_z"
## [34] "accel_dumbbell_x"    "accel_dumbbell_y"    "accel_dumbbell_z"
## [37] "magnet_dumbbell_x"   "magnet_dumbbell_y"   "magnet_dumbbell_z"
## [40] "roll_forearm"        "pitch_forearm"       "yaw_forearm"

```

```
## [43] "total_accel_forearm" "gyros_forearm_x" "gyros_forearm_y"
## [46] "gyros_forearm_z" "accel_forearm_x" "accel_forearm_y"
## [49] "accel_forearm_z" "magnet_forearm_x" "magnet_forearm_y"
## [52] "magnet_forearm_z" "classe"
```

```
colnames(df_testing)
```

```
## [1] "roll_belt" "pitch_belt" "yaw_belt"
## [4] "total_accel_belt" "gyros_belt_x" "gyros_belt_y"
## [7] "gyros_belt_z" "accel_belt_x" "accel_belt_y"
## [10] "accel_belt_z" "magnet_belt_x" "magnet_belt_y"
## [13] "magnet_belt_z" "roll_arm" "pitch_arm"
## [16] "yaw_arm" "total_accel_arm" "gyros_arm_x"
## [19] "gyros_arm_y" "gyros_arm_z" "accel_arm_x"
## [22] "accel_arm_y" "accel_arm_z" "magnet_arm_x"
## [25] "magnet_arm_y" "magnet_arm_z" "roll_dumbbell"
## [28] "pitch_dumbbell" "yaw_dumbbell" "total_accel_dumbbell"
## [31] "gyros_dumbbell_x" "gyros_dumbbell_y" "gyros_dumbbell_z"
## [34] "accel_dumbbell_x" "accel_dumbbell_y" "accel_dumbbell_z"
## [37] "magnet_dumbbell_x" "magnet_dumbbell_y" "magnet_dumbbell_z"
## [40] "roll_forearm" "pitch_forearm" "yaw_forearm"
## [43] "total_accel_forearm" "gyros_forearm_x" "gyros_forearm_y"
## [46] "gyros_forearm_z" "accel_forearm_x" "accel_forearm_y"
## [49] "accel_forearm_z" "magnet_forearm_x" "magnet_forearm_y"
## [52] "magnet_forearm_z" "problem_id"
```

First, check for covariates that have virtually no variability.

```
nsv <- nearZeroVar(df_training, saveMetrics=TRUE)
nsv
```

```
##          freqRatio percentUnique zeroVar  nsv
## roll_belt      1.101904      6.7781062 FALSE FALSE
## pitch_belt      1.036082      9.3772296 FALSE FALSE
## yaw_belt        1.058480      9.9734991 FALSE FALSE
## total_accel_belt 1.063160      0.1477933 FALSE FALSE
## gyros_belt_x     1.058651      0.7134849 FALSE FALSE
## gyros_belt_y     1.144000      0.3516461 FALSE FALSE
## gyros_belt_z     1.066214      0.8612782 FALSE FALSE
## accel_belt_x     1.055412      0.8357966 FALSE FALSE
## accel_belt_y     1.113725      0.7287738 FALSE FALSE
## accel_belt_z     1.078767      1.5237998 FALSE FALSE
## magnet_belt_x    1.090141      1.6664968 FALSE FALSE
## magnet_belt_y    1.099688      1.5187035 FALSE FALSE
## magnet_belt_z    1.006369      2.3290184 FALSE FALSE
## roll_arm        52.338462     13.5256345 FALSE FALSE
## pitch_arm       87.256410     15.7323412 FALSE FALSE
## yaw_arm         33.029126     14.6570176 FALSE FALSE
## total_accel_arm  1.024526      0.3363572 FALSE FALSE
## gyros_arm_x      1.015504      3.2769341 FALSE FALSE
## gyros_arm_y      1.454369      1.9162165 FALSE FALSE
```

## gyros_arm_z	1.110687	1.2638875	FALSE	FALSE
## accel_arm_x	1.017341	3.9598410	FALSE	FALSE
## accel_arm_y	1.140187	2.7367241	FALSE	FALSE
## accel_arm_z	1.128000	4.0362858	FALSE	FALSE
## magnet_arm_x	1.000000	6.8239731	FALSE	FALSE
## magnet_arm_y	1.056818	4.4439914	FALSE	FALSE
## magnet_arm_z	1.036364	6.4468454	FALSE	FALSE
## roll_dumbbell	1.022388	84.2065029	FALSE	FALSE
## pitch_dumbbell	2.277372	81.7449801	FALSE	FALSE
## yaw_dumbbell	1.132231	83.4828254	FALSE	FALSE
## total_accel_dumbbell	1.072634	0.2191418	FALSE	FALSE
## gyros_dumbbell_x	1.003268	1.2282132	FALSE	FALSE
## gyros_dumbbell_y	1.264957	1.4167771	FALSE	FALSE
## gyros_dumbbell_z	1.060100	1.0498420	FALSE	FALSE
## accel_dumbbell_x	1.018018	2.1659362	FALSE	FALSE
## accel_dumbbell_y	1.053061	2.3748853	FALSE	FALSE
## accel_dumbbell_z	1.133333	2.0894914	FALSE	FALSE
## magnet_dumbbell_x	1.098266	5.7486495	FALSE	FALSE
## magnet_dumbbell_y	1.197740	4.3012945	FALSE	FALSE
## magnet_dumbbell_z	1.020833	3.4451126	FALSE	FALSE
## roll_forearm	11.589286	11.0895933	FALSE	FALSE
## pitch_forearm	65.983051	14.8557741	FALSE	FALSE
## yaw_forearm	15.322835	10.1467740	FALSE	FALSE
## total_accel_forearm	1.128928	0.3567424	FALSE	FALSE
## gyros_forearm_x	1.059273	1.5187035	FALSE	FALSE
## gyros_forearm_y	1.036554	3.7763735	FALSE	FALSE
## gyros_forearm_z	1.122917	1.5645704	FALSE	FALSE
## accel_forearm_x	1.126437	4.0464784	FALSE	FALSE
## accel_forearm_y	1.059406	5.1116094	FALSE	FALSE
## accel_forearm_z	1.006250	2.9558659	FALSE	FALSE
## magnet_forearm_x	1.012346	7.7667924	FALSE	FALSE
## magnet_forearm_y	1.246914	9.5403119	FALSE	FALSE
## magnet_forearm_z	1.000000	8.5771073	FALSE	FALSE
## classe	1.469581	0.0254816	FALSE	FALSE

Given that all of the near zero variance variables (nsv) are FALSE, there's no need to eliminate any covariates due to lack of variability.

## ALGORITHM

We were provided with a large training set (19,622 entries) and a small testing set (20 entries). Instead of performing the algorithm on the entire training set, as it would be time consuming and wouldn't allow for an attempt on a testing set, I chose to divide the given training set into four roughly equal sets, each of which was then split into a training set (comprising 60% of the entries) and a testing set (comprising 40% of the entries).

```
# Divide the given training set into 4 roughly equal sets.
set.seed(666)
ids_small <- createDataPartition(y=df_training$classe, p=0.25, list=FALSE)
df_small1 <- df_training[ids_small,]
```

```

df_remainder <- df_training[-ids_small,]
set.seed(666)
ids_small <- createDataPartition(y=df_remainder$classe, p=0.33, list=FALSE)
df_small12 <- df_remainder[ids_small,]
df_remainder <- df_remainder[-ids_small,]
set.seed(666)
ids_small <- createDataPartition(y=df_remainder$classe, p=0.5, list=FALSE)
df_small13 <- df_remainder[ids_small,]
df_small14 <- df_remainder[-ids_small,]
# Divide each of these 4 sets into training (60%) and test (40%) sets.
set.seed(666)
inTrain <- createDataPartition(y=df_small11$classe, p=0.6, list=FALSE)
df_small_training1 <- df_small11[inTrain,]
df_small_testing1 <- df_small11[-inTrain,]
set.seed(666)
inTrain <- createDataPartition(y=df_small12$classe, p=0.6, list=FALSE)
df_small_training2 <- df_small12[inTrain,]
df_small_testing2 <- df_small12[-inTrain,]
set.seed(666)
inTrain <- createDataPartition(y=df_small13$classe, p=0.6, list=FALSE)
df_small_training3 <- df_small13[inTrain,]
df_small_testing3 <- df_small13[-inTrain,]
set.seed(666)
inTrain <- createDataPartition(y=df_small14$classe, p=0.6, list=FALSE)
df_small_training4 <- df_small14[inTrain,]
df_small_testing4 <- df_small14[-inTrain,]

```

## PARAMETERS

I decided to try classification trees "out of the box" and then introduce preprocessing and cross validation.

While I also considered applying "out of the box" random forest models, some of the horror stories contributed to the coursera discussion forums regarding the lengthy processing times for random forest models convinced me to only attempt random forests with cross validation and, possibly, preprocessing.

## EVALUATION

Classification Tree

First, the "out of the box" classification tree:

```

# Train on training set 1 of 4 with no extra features.
set.seed(666)
modFit <- train(df_small_training1$classe ~ ., data = df_small_training1,
method="rpart")
print(modFit, digits=3)

```

```

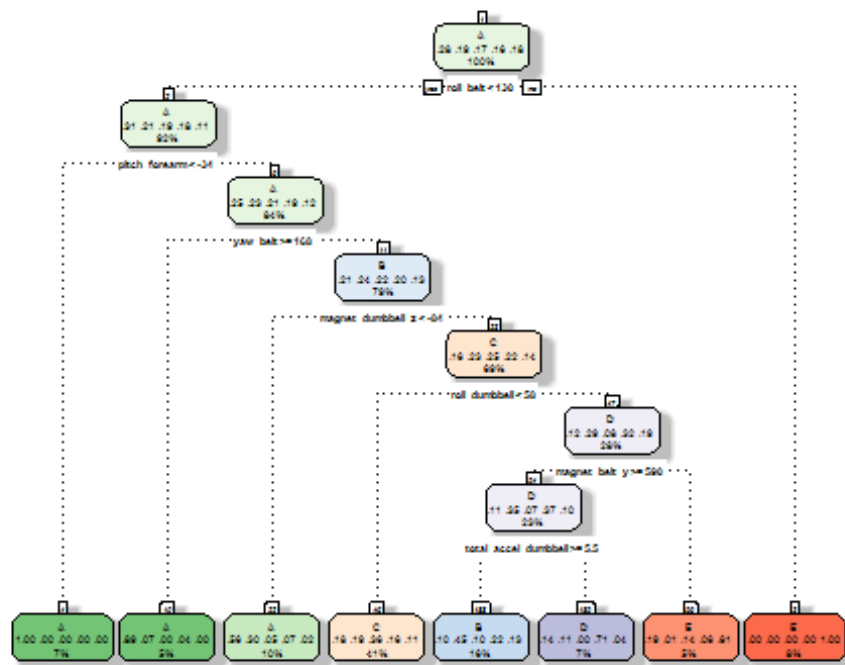
## CART
##
## 2946 samples
## 52 predictor
## 5 classes: 'A', 'B', 'C', 'D', 'E'
##
## No pre-processing
## Resampling: Bootstrapped (25 reps)
## Summary of sample sizes: 2946, 2946, 2946, 2946, 2946, 2946, ...
## Resampling results across tuning parameters:
##
##   cp      Accuracy   Kappa   Accuracy SD   Kappa SD
##   0.0346  0.531      0.4003  0.0355        0.0479
##   0.0442  0.471      0.3076  0.0555        0.0967
##   0.1162  0.324      0.0602  0.0456        0.0641
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was cp = 0.0346.

print(modFit$finalModel, digits=3)

## n= 2946
##
## node), split, n, loss, yval, (yprob)
##      * denotes terminal node
##
## 1) root 2946 2110 A (0.28 0.19 0.17 0.16 0.18)
##    2) roll_belt< 130 2699 1860 A (0.31 0.21 0.19 0.18 0.11)
##      4) pitch_forearm< -34 220 0 A (1 0 0 0 0) *
##      5) pitch_forearm>=-34 2479 1860 A (0.25 0.23 0.21 0.19 0.12)
##        10) yaw_belt>=168 138 15 A (0.89 0.072 0 0.036 0) *
##        11) yaw_belt< 168 2341 1780 B (0.21 0.24 0.22 0.2 0.13)
##          22) magnet_dumbbell_z< -83.5 305 134 A (0.56 0.3 0.046 0.069
0.02) *
##            23) magnet_dumbbell_z>=-83.5 2036 1540 C (0.16 0.23 0.25 0.22
0.14)
##              46) roll_dumbbell< 57.7 1209 776 C (0.18 0.19 0.36 0.16 0.11)
*
##                47) roll_dumbbell>=57.7 827 565 D (0.12 0.29 0.081 0.32 0.19)
##                  94) magnet_belt_y>=590 687 433 D (0.11 0.35 0.07 0.37 0.1)
##                    188) total_accel_dumbbell>=5.5 474 260 B (0.097 0.45 0.1
0.22 0.13) *
##                      189) total_accel_dumbbell< 5.5 213 62 D (0.14 0.11 0 0.71
0.042) *
##                        95) magnet_belt_y< 590 140 55 E (0.19 0.014 0.14 0.057
0.61) *
##          3) roll_belt>=130 247 1 E (0.004 0 0 0 1) *

fancyRpartPlot(modFit$finalModel)

```



```
# Run against testing set 1 of 4 with no extra features.
```

```
predictions <- predict(modFit, newdata=df_small_testing1)
```

```
print(confusionMatrix(predictions, df_small_testing1$classe), digits=4)
```

## ## Confusion Matrix and Statistics

##

## ## Reference

## ## Prediction

```
##          A 368   74   11   28    8
```

##	B	24	151	25	83	30
----	---	----	-----	----	----	----

```
##          C 135 148 288 138  99
```

```
##          D   15    7    0   69    4
```

```
##          E  16   0  18   3 219
```

### ##

##

### ##

##

```
##          Accrual day : 0.5361
##          95% CI : (0.5361,
```

```
##          95% CI : (0.9999, 0.9999)
##      No Information Rate : 0.2845
```

```
## P-Value [Acc > NIR] : < 2.2e
```

```
##      +-----+-----+-----+-----+-----+-----+-----+-----+
##
```

##

```
## McNemar's Test P-Value : < 2.2e
```

```
##
##
```

##

```
## Statistics by class.
##
```

##

```
##          class. A class. B class. C class. D class. E
```



## Sensitivity	0.6595	0.3974	0.8421	0.21495	0.6083
## Specificity	0.9138	0.8975	0.6788	0.98415	0.9769
## Pos Pred Value	0.7526	0.4824	0.3564	0.72632	0.8555
## Neg Pred Value	0.8709	0.8610	0.9532	0.86495	0.9173
## Prevalence	0.2845	0.1938	0.1744	0.16369	0.1836
## Detection Rate	0.1877	0.0770	0.1469	0.03519	0.1117
## Detection Prevalence	0.2494	0.1596	0.4120	0.04844	0.1305
## Balanced Accuracy	0.7866	0.6475	0.7605	0.59955	0.7926

low accuracy rate (0.5584)

*# Train on training set 1 of 4 with only preprocessing.*

```
set.seed(666)
modFit <- train(df_small_training1$classe ~ ., preProcess=c("center",
"scale"), data = df_small_training1, method="rpart")
print(modFit, digits=3)
```

```
## CART
##
## 2946 samples
## 52 predictor
## 5 classes: 'A', 'B', 'C', 'D', 'E'
##
## Pre-processing: centered, scaled
## Resampling: Bootstrapped (25 reps)
## Summary of sample sizes: 2946, 2946, 2946, 2946, 2946, 2946, ...
## Resampling results across tuning parameters:
##
## cp      Accuracy  Kappa  Accuracy SD  Kappa SD
## 0.0346  0.531      0.4003  0.0355      0.0479
## 0.0442  0.471      0.3077  0.0555      0.0968
## 0.1162  0.324      0.0602  0.0456      0.0641
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was cp = 0.0346.
```

*# Train on training set 1 of 4 with only cross validation.*

```
set.seed(666)
modFit <- train(df_small_training1$classe ~ ., trControl=trainControl(method
= "cv", number = 4), data = df_small_training1, method="rpart")
print(modFit, digits=3)
```

```
## CART
##
## 2946 samples
## 52 predictor
## 5 classes: 'A', 'B', 'C', 'D', 'E'
##
## No pre-processing
## Resampling: Cross-Validated (4 fold)
## Summary of sample sizes: 2212, 2209, 2208, 2209
```

```

## Resampling results across tuning parameters:
##
##   cp      Accuracy  Kappa  Accuracy SD  Kappa SD
##   0.0346  0.552     0.4266  0.0383      0.0542
##   0.0442  0.470     0.3041  0.0689      0.1197
##   0.1162  0.344     0.0914  0.0405      0.0610
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was cp = 0.0346.

# Train on training set 1 of 4 with both preprocessing and cross validation.
set.seed(666)
modFit <- train(df_small_training1$classe ~ ., preProcess=c("center",
"scale"), trControl=trainControl(method = "cv", number = 4), data =
df_small_training1, method="rpart")
print(modFit, digits=3)

## CART
##
## 2946 samples
## 52 predictor
## 5 classes: 'A', 'B', 'C', 'D', 'E'
##
## Pre-processing: centered, scaled
## Resampling: Cross-Validated (4 fold)
## Summary of sample sizes: 2212, 2209, 2208, 2209
## Resampling results across tuning parameters:
##
##   cp      Accuracy  Kappa  Accuracy SD  Kappa SD
##   0.0346  0.552     0.4266  0.0383      0.0542
##   0.0442  0.470     0.3041  0.0689      0.1197
##   0.1162  0.344     0.0914  0.0405      0.0610
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was cp = 0.0346.

# Run against testing set 1 of 4 with both preprocessing and cross
validation.
predictions <- predict(modFit, newdata=df_small_testing1)
print(confusionMatrix(predictions, df_small_testing1$classe), digits=4)

## Confusion Matrix and Statistics
##
##           Reference
## Prediction  A   B   C   D   E
##      A 368  74  11  28   8
##      B  24 151  25  83  30
##      C 135 148 288 138  99
##      D  15   7   0  69   4
##      E  16   0  18   3 219
##

```

```
## Overall Statistics
##
##           Accuracy : 0.5584
##           95% CI : (0.5361, 0.5805)
##      No Information Rate : 0.2845
##      P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.4441
##  McNemar's Test P-Value : < 2.2e-16
##
## Statistics by Class:
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity      0.6595  0.3974  0.8421  0.21495  0.6083
## Specificity      0.9138  0.8975  0.6788  0.98415  0.9769
## Pos Pred Value   0.7526  0.4824  0.3564  0.72632  0.8555
## Neg Pred Value   0.8709  0.8610  0.9532  0.86495  0.9173
## Prevalence       0.2845  0.1938  0.1744  0.16369  0.1836
## Detection Rate   0.1877  0.0770  0.1469  0.03519  0.1117
## Detection Prevalence 0.2494  0.1596  0.4120  0.04844  0.1305
## Balanced Accuracy 0.7866  0.6475  0.7605  0.59955  0.7926
```

The impact of incorporating both preprocessing and cross validation appeared to show some minimal improvement (accuracy rate rose from 0.531 to 0.552 against training sets). However, when run against the corresponding testing set, the accuracy rate was identical (0.5584) for both the "out of the box" and the preprocessing/cross validation methods.

## Random Forest

First I decided to assess the impact/value of including preprocessing.

```
# Train on training set 1 of 4 with only cross validation.
set.seed(666)
modFit <- train(df_small_training1$classe ~ ., method="rf",
trControl=trainControl(method = "cv", number = 4), data=df_small_training1)
print(modFit, digits=3)

## Random Forest
##
## 2946 samples
## 52 predictor
## 5 classes: 'A', 'B', 'C', 'D', 'E'
##
## No pre-processing
## Resampling: Cross-Validated (4 fold)
## Summary of sample sizes: 2212, 2209, 2208, 2209
## Resampling results across tuning parameters:
##
## mtry Accuracy Kappa Accuracy SD Kappa SD
## 2 0.951 0.939 0.00291 0.00367
```

```
## 27 0.957 0.945 0.00740 0.00937
## 52 0.951 0.939 0.01226 0.01549
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was mtry = 27.

# Run against testing set 1 of 4.
predictions <- predict(modFit, newdata=df_small_testing1)
print(confusionMatrix(predictions, df_small_testing1$classe), digits=4)

## Confusion Matrix and Statistics
##
##           Reference
## Prediction  A   B   C   D   E
##           A 556  10   1   0   0
##           B   2 360  13   0   2
##           C   0   9 323   4   5
##           D   0   1   5 313   2
##           E   0   0   0   4 351
##
## Overall Statistics
##
##           Accuracy : 0.9704
##           95% CI : (0.9619, 0.9775)
##           No Information Rate : 0.2845
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.9626
##           McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity      0.9964  0.9474  0.9444  0.9751  0.9750
## Specificity      0.9922  0.9892  0.9889  0.9951  0.9975
## Pos Pred Value   0.9806  0.9549  0.9472  0.9751  0.9887
## Neg Pred Value   0.9986  0.9874  0.9883  0.9951  0.9944
## Prevalence       0.2845  0.1938  0.1744  0.1637  0.1836
## Detection Rate   0.2835  0.1836  0.1647  0.1596  0.1790
## Detection Prevalence 0.2891  0.1922  0.1739  0.1637  0.1810
## Balanced Accuracy 0.9943  0.9683  0.9667  0.9851  0.9863

# Run against 20 testing set
print(predict(modFit, newdata=df_testing))

## [1] B A A A A E D B A A B C B A E E A B B B
## Levels: A B C D E

# Train on training set 1 of 4 with only both preprocessing and cross
validation.
set.seed(666)
```

```

modFit <- train(df_small_training1$classe ~ ., method="rf",
preProcess=c("center", "scale"), trControl=trainControl(method = "cv", number
= 4), data=df_small_training1)
print(modFit, digits=3)

## Random Forest
##
## 2946 samples
## 52 predictor
## 5 classes: 'A', 'B', 'C', 'D', 'E'
##
## Pre-processing: centered, scaled
## Resampling: Cross-Validated (4 fold)
## Summary of sample sizes: 2212, 2209, 2208, 2209
## Resampling results across tuning parameters:
##
## mtry Accuracy Kappa Accuracy SD Kappa SD
## 2 0.951 0.939 0.00172 0.00217
## 27 0.955 0.943 0.00588 0.00743
## 52 0.952 0.939 0.01061 0.01341
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was mtry = 27.

# Run against testing set 1 of 4.
predictions <- predict(modFit, newdata=df_small_testing1)
print(confusionMatrix(predictions, df_small_testing1$classe), digits=4)

## Confusion Matrix and Statistics
##
## Reference
## Prediction A B C D E
## A 556 12 1 0 1
## B 2 358 12 0 0
## C 0 9 324 6 5
## D 0 1 5 310 2
## E 0 0 0 5 352
##
## Overall Statistics
##
## Accuracy : 0.9689
## 95% CI : (0.9602, 0.9761)
## No Information Rate : 0.2845
## P-Value [Acc > NIR] : < 2.2e-16
##
## Kappa : 0.9606
## McNemar's Test P-Value : NA
##
## Statistics by Class:
##
## Class: A Class: B Class: C Class: D Class: E

```

```
## Sensitivity          0.9964    0.9421    0.9474    0.9657    0.9778
## Specificity          0.9900    0.9911    0.9876    0.9951    0.9969
## Pos Pred Value       0.9754    0.9624    0.9419    0.9748    0.9860
## Neg Pred Value       0.9986    0.9862    0.9889    0.9933    0.9950
## Prevalence           0.2845    0.1938    0.1744    0.1637    0.1836
## Detection Rate       0.2835    0.1826    0.1652    0.1581    0.1795
## Detection Prevalence 0.2907    0.1897    0.1754    0.1622    0.1820
## Balanced Accuracy     0.9932    0.9666    0.9675    0.9804    0.9873
```

*# Run against 20 testing*

```
print(predict(modFit, newdata=df_testing))
```

```
## [1] B A A A A E D B A A B C B A E E A B B B
```

```
## Levels: A B C D E
```

Preprocessing actually lowered the accuracy rate from 0.955 to 0.954 against the training set. However, when run against the corresponding set, the accuracy rate rose from 0.9689 to 0.9714 with the addition of preprocessing. Thus I decided to apply both preprocessing and cross validation to the remaining 3 data sets.

*# Train on training set 2 of 4 with only cross validation.*

```
set.seed(666)
```

```
modFit <- train(df_small_training2$classe ~ ., method="rf",
preProcess=c("center", "scale"), trControl=trainControl(method = "cv", number
= 4), data=df_small_training2)
```

```
print(modFit, digits=3)
```

```
## Random Forest
```

```
##
```

```
## 2917 samples
```

```
## 52 predictor
```

```
## 5 classes: 'A', 'B', 'C', 'D', 'E'
```

```
##
```

```
## Pre-processing: centered, scaled
```

```
## Resampling: Cross-Validated (4 fold)
```

```
## Summary of sample sizes: 2188, 2188, 2187, 2188
```

```
## Resampling results across tuning parameters:
```

```
##
```

```
## mtry Accuracy Kappa Accuracy SD Kappa SD
```

```
## 2 0.953 0.941 0.00953 0.01210
```

```
## 27 0.952 0.939 0.00699 0.00889
```

```
## 52 0.941 0.926 0.00539 0.00683
```

```
##
```

```
## Accuracy was used to select the optimal model using the largest value.
```

```
## The final value used for the model was mtry = 2.
```

*# Run against testing set 2 of 4.*

```
predictions <- predict(modFit, newdata=df_small_testing2)
```

```
print(confusionMatrix(predictions, df_small_testing2$classe), digits=4)
```

```
## Confusion Matrix and Statistics
```

```
##
```

```

##           Reference
## Prediction   A    B    C    D    E
##           A 547  12    0    3    0
##           B   2 351  22    0    1
##           C   0  12 314   19    6
##           D   2   1   2 293    6
##           E   1   0   0   3 344
##
## Overall Statistics
##
##           Accuracy : 0.9526
##           95% CI : (0.9422, 0.9616)
##           No Information Rate : 0.2844
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.94
## Mcnemar's Test P-Value : NA
##
## Statistics by Class:
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity      0.9909  0.9335  0.9290  0.9214  0.9636
## Specificity      0.9892  0.9840  0.9769  0.9932  0.9975
## Pos Pred Value   0.9733  0.9335  0.8946  0.9638  0.9885
## Neg Pred Value   0.9964  0.9840  0.9849  0.9847  0.9918
## Prevalence       0.2844  0.1937  0.1741  0.1638  0.1839
## Detection Rate   0.2818  0.1808  0.1618  0.1510  0.1772
## Detection Prevalence 0.2895  0.1937  0.1808  0.1566  0.1793
## Balanced Accuracy 0.9901  0.9588  0.9530  0.9573  0.9805

# Run against 20 testing set
print(predict(modFit, newdata=df_testing))

## [1] B A B A A E D B A A B C B A E E A B B B
## Levels: A B C D E

# Train on training set 3 of 4 with only cross validation.
set.seed(666)
modFit <- train(df_small_training3$classe ~ ., method="rf",
preProcess=c("center", "scale"), trControl=trainControl(method = "cv", number
= 4), data=df_small_training3)
print(modFit, digits=3)

## Random Forest
##
## 2960 samples
## 52 predictor
## 5 classes: 'A', 'B', 'C', 'D', 'E'
##
## Pre-processing: centered, scaled
## Resampling: Cross-Validated (4 fold)

```

```

## Summary of sample sizes: 2219, 2221, 2220, 2220
## Resampling results across tuning parameters:
##
##   mtry Accuracy  Kappa Accuracy SD  Kappa SD
##    2   0.949    0.936  0.00572    0.00724
##   27   0.949    0.936  0.00953    0.01206
##   52   0.942    0.926  0.01020    0.01293
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was mtry = 27.

# Run against testing set 3 of 4.
predictions <- predict(modFit, newdata=df_small_testing3)
print(confusionMatrix(predictions, df_small_testing3$classe), digits=4)

## Confusion Matrix and Statistics
##
##           Reference
## Prediction  A    B    C    D    E
##           A 555  10    0    1    0
##           B   1 358  18    0    3
##           C   1  12 320    7    3
##           D   2   1   3 313    1
##           E   1   0   3   2 355
##
## Overall Statistics
##
##               Accuracy : 0.965
##               95% CI : (0.9559, 0.9726)
##       No Information Rate : 0.2843
##       P-Value [Acc > NIR] : <2e-16
##
##               Kappa : 0.9557
##  McNemar's Test P-Value : 0.0782
##
## Statistics by Class:
##
##               Class: A Class: B Class: C Class: D Class: E
## Sensitivity      0.9911  0.9396  0.9302  0.9690  0.9807
## Specificity      0.9922  0.9862  0.9859  0.9957  0.9963
## Pos Pred Value   0.9806  0.9421  0.9329  0.9781  0.9834
## Neg Pred Value   0.9964  0.9855  0.9852  0.9939  0.9956
## Prevalence       0.2843  0.1934  0.1746  0.1640  0.1838
## Detection Rate   0.2817  0.1817  0.1624  0.1589  0.1802
## Detection Prevalence 0.2873  0.1929  0.1741  0.1624  0.1832
## Balanced Accuracy 0.9916  0.9629  0.9580  0.9824  0.9885

# Run against 20 testing set
print(predict(modFit, newdata=df_testing))

```



```
## [1] B A B A A E D B A A B C B A E E A B B B
## Levels: A B C D E

# Train on training set 4 of 4 with only cross validation.
set.seed(666)
modFit <- train(df_small_training4$classe ~ ., method="rf",
preProcess=c("center", "scale"), trControl=trainControl(method = "cv", number
= 4), data=df_small_training4)
print(modFit, digits=3)

## Random Forest
##
## 2958 samples
## 52 predictor
## 5 classes: 'A', 'B', 'C', 'D', 'E'
##
## Pre-processing: centered, scaled
## Resampling: Cross-Validated (4 fold)
## Summary of sample sizes: 2218, 2219, 2219, 2218
## Resampling results across tuning parameters:
##
## mtry Accuracy Kappa Accuracy SD Kappa SD
## 2 0.947 0.933 0.00969 0.01228
## 27 0.957 0.945 0.00722 0.00914
## 52 0.947 0.933 0.01031 0.01307
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was mtry = 27.

# Run against testing set 4 of 4.
predictions <- predict(modFit, newdata=df_small_testing4)
print(confusionMatrix(predictions, df_small_testing4$classe), digits=4)

## Confusion Matrix and Statistics
##
## Reference
## Prediction A B C D E
## A 552 19 0 0 0
## B 5 358 20 3 2
## C 2 4 315 8 7
## D 1 0 8 311 6
## E 0 0 0 1 347
##
## Overall Statistics
##
## Accuracy : 0.9563
## 95% CI : (0.9463, 0.9649)
## No Information Rate : 0.2844
## P-Value [Acc > NIR] : < 2.2e-16
##
## Kappa : 0.9447
```

```
## McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity      0.9857  0.9396  0.9184  0.9628  0.9586
## Specificity      0.9865  0.9811  0.9871  0.9909  0.9994
## Pos Pred Value   0.9667  0.9227  0.9375  0.9540  0.9971
## Neg Pred Value   0.9943  0.9855  0.9829  0.9927  0.9907
## Prevalence       0.2844  0.1935  0.1742  0.1640  0.1838
## Detection Rate   0.2803  0.1818  0.1600  0.1579  0.1762
## Detection Prevalence 0.2900  0.1971  0.1706  0.1656  0.1767
## Balanced Accuracy 0.9861  0.9604  0.9527  0.9769  0.9790

# Run against 20 testing set provided by Professor Leek.
print(predict(modFit, newdata=df_testing))

## [1] B A B A A E D D A A B C B A E E A B B B
## Levels: A B C D E
```

## CONCLUSION

I received three separate predictions by applying the 4 models against the actual 20 item training set:

- A) Accuracy Rate 0.0286 Predictions: B A A A A E D B A A B C B A E E A B B B
- B) Accuracy Rates 0.0366 and 0.0345 Predictions: B A B A A E D B A A B C B A E E A B B B
- C) Accuracy Rate 0.0437 Predictions: B A B A A E D D A A B C B A E E A B B B

Since Professor Leek is allowing 2 submissions for each problem, I decided to attempt with the two most likely prediction sets: option A and option B.

Since options A and B above only differed for item 3 (A for option A, B for option B), I submitted one value for problems 1-2 and 4-20, while I submitted two values for problem 3. For problem 3, I was expecting the automated grader to tell me which answer (A or B) was correct, but instead the grader simply told me I had a correct answer. All other answers were also correct, resulting in a score of 100%.